

# HW7\_edgarsp2

November 25, 2019

## 1 STAT 542 / CS 598: Homework 7

Fall 2019, by Edgar Pino (edgarsp2)

Due: Monday, Nov 25 by 11:59 PM Pacific Time

```
[62]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
```

## 2 Question 1 [100 Points] AdaBoost with stump model

### 2.1 Generate Data

```
[63]: def plot_data(x, y):
    unique = np.unique(y)
    y_plot = y + 0.1*np.random.uniform(-1, 1, len(x))
    for li in range(len(unique)):
        _x = x[y == unique[li]]
        _y = y_plot[y == unique[li]]
        label = 'Positive' if unique[li] == 1 else 'Negative'
        plt.scatter(_x, _y, c = COLOR_LABELS[li], label=label, s=8)

    # lines(sort(x), py(x)[order(x)] - 0.5)
    # plt.plot(x, py(y)-0.5, c='blue')
    plt.legend()
    plt.show()
```

```
[64]: np.random.seed(1)
```

```
[65]: COLOR_LABELS = ["deepskyblue", "darkorange"]
```


```
[66]: n = 300
```

```
[67]: x = np.random.uniform(size=n)
```

```
[68]: py = lambda x: np.sin(4*np.pi*x)/3 + 0.5
```

```
[69]: y = (np.random.binomial(1, py(x), n)-0.5)*2
```

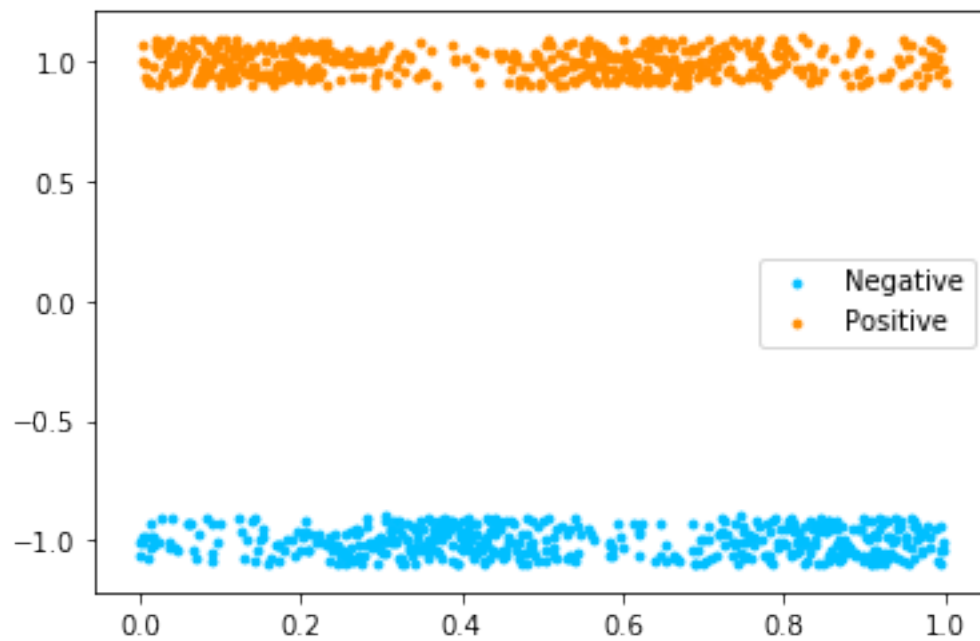
```
[70]: plot_data(x,y)
```



```
[71]: x_test = np.linspace(0, 1, 1000)
```

```
[72]: y_test = (np.random.binomial(1, py(x_test), 1000)-0.5)*2
```

```
[73]: plot_data(x_test,y_test)
```



## 2.2 Create model

```
[74]: def exponentialLoss(y, ypred):  
       return np.mean(np.exp(-(2. * y - 1.) * ypred))  
  
[75]: class Stump():  
       def __init__(self):  
           self.label = 1  
           self.threshold = None  
           self.loss = None  
           self.alpha = None  
  
[76]: class MyAdaboost():  
       def __init__(self, shrinkage=.5, num_classifiers=10):  
           self.num_classifiers = num_classifiers  
           self.shrinkage = shrinkage  
           self.weights = None  
  
       def fit(self, x, y):  
           n_samples = len(x)  
           weights = np.full(n_samples, (1 / n_samples))  
  
           self.classifiers = []  
  
           for _ in range(self.num_classifiers):  
               classifier = Stump()  
               min_error = float('inf')  
  
               feature_values = np.expand_dims(x, axis=1)  
               unique_values = np.unique(feature_values)  
  
               for threshold in unique_values:  
                   p = 1  
                   prediction = np.ones(np.shape(y))  
                   prediction[x < threshold] = -1  
                   error = sum(weights[y != prediction])  
  
                   if error > 0.5:  
                       error = 1 - error  
                       p = -1  
  
                   if error < min_error:  
                       classifier.label = p  
                       classifier.threshold = threshold  
                       min_error = error  
                       classifier.loss = exponentialLoss(y, prediction)
```

```

        classifier.alpha = self.shrinkage * np.log((1.0 - min_error) /
→(min_error + 1e-10))
        predictions = np.ones(np.shape(y))
        negative_idx = (classifier.label * x < classifier.label *
→classifier.threshold)
        predictions[negative_idx] = -1
        weights *= np.exp(-classifier.alpha * y * predictions)
        weights /= np.sum(weights)

        self.classifiers.append(classifier)

    self.weights = weights

    def predict(self, x):
        n_samples = len(x)
        y_pred = np.zeros((n_samples, 1))

        for classifier in self.classifiers:
            predictions = np.ones(np.shape(y_pred))
            negative_idx = (classifier.label * x < classifier.label *
→classifier.threshold)
            predictions[negative_idx] = -1

            y_pred += classifier.alpha * predictions
        return np.sign(y_pred).flatten()

```

```
[77]: number_of_classifiers = 150
```

```
[78]: model = MyAdaboost(shrinkage=.5, num_classifiers=number_of_classifiers)
```

```
[79]: model.fit(x, y)
```

```
[80]: y_pred = model.predict(x_test)
```


```
[81]: f"Model accuracy {accuracy_score(y_test, y_pred) * 100}%"
```

```
[81]: 'Model accuracy 67.30000000000001%'
```

## 2.3 Plot the exponential loss

```
[82]: loss_values = [clf.loss for clf in model.classifiers]
```


```
[83]: plt.plot(loss_values)
plt.show()
```



It appears that the best number of classifiers is about 120. After that, the model loss seems to stay the same.

## 2.4 Try a few different shrinkage factors and comment on your findings

```
[84]: shrinkage_values = np.arange(.1,1,.1)
[85]: scores = []
[86]: for val in shrinkage_values:
    s_model = MyAdaboost(shrinkage=val, num_classifiers=number_of_classifiers)
    s_model.fit(x, y)
    y_pred = s_model.predict(x_test)
    score = accuracy_score(y_test, y_pred) * 100
    scores.append(score)
[87]: plt.plot(shrinkage_values, scores)
plt.show()
```



Based on this chart, the best shrinkage value is between .1 and .3. After that, the model accuracy score drops.

## 2.5 Plot the final model

```
[88]: best_num_classifiers=120
```

```
[89]: best_shrinkage=.1
```

```
[90]: final_model = MyAdaboost(shrinkage=best_shrinkage,
    →num_classifiers=best_num_classifiers)
```


```
[91]: final_model.fit(x, y)
```

```
[92]: y_pred = final_model.predict(x_test)
```

```
[93]: f"Final Model accuracy {accuracy_score(y_test, y_pred) * 100}"
```


```
[93]: 'Final Model accuracy 68.8'
```

```
[94]: plt.plot(final_model.weights)
plt.show()
```



```
[ ]: plt.plot(final_model.weights)
plt.show()
```

```
[61]: plot_data(x_test,y_pred)
```



[: