

# HW1\_edgarsp2

September 9, 2019

## 1 STAT 542 / CS 598: Example Homework

Spring 2019, by Edgar Pino (edgarsp2)

Due: Monday, Sep 9 by 11:59 PM Pacific Time.

### 1.1 Question 1 [50 Points] KNN

1A. [15 Points] Write a function `myknn(xtest, xtrain, ytrain, k)` that fits a KNN model that predict a target point or multiple target points `xtest`. Here `xtrain` is the training dataset covariate value, `ytrain` is the training data outcome, and `k` is the number of nearest neighbors. Use the 2 norm to evaluate the distance between two points. Please note that you cannot use any additional R package within this function.

#### 1.1.1 Answer:

```
import operator
import pandas as pd
import numpy as np

def get_distance(a, b):
    a = np.array(a)
    b = np.array(b)
    return np.sqrt(np.sum(np.square(a - b)))

def get_label_average(neighbors):
    return np.average(neighbors[:, -1])

def get_neighbors(trainingSet, instance, k):
    distances = []

    for x in range(len(trainingSet)):
        distance = get_distance(instance, trainingSet[x] [: -1])
        distances.append((trainingSet[x], distance))

    distances.sort(key=operator.itemgetter(1))

    neighbors = []
    for x in range(k):
```

```

        neighbors.append(distances[x][0])
    return np.array(neighbors)

def myknn(xtest, xtrain, ytrain, k):
    nn = []
    training_data_xy = []

    for i in range(len(xtrain)):
        training_data_xy.append(np.append(xtrain[i], [ytrain[i]]))

    for i in range(len(xtest)):
        neighbors = get_neighbors(training_data_xy, xtest[i], k)
        predicted_label = get_label_average(neighbors)
        nn.append(predicted_label)

    return np.array(nn)

```

## 1.2 1B. [10 Points] Generate 1000 observations from a five-dimensional normally distribution.

### 1.2.1 Answer:

```
np.random.seed(1)
```

```

def get_cov():
    cov = np.zeros((5,5))
    for i in range(1,6):
        for j in range(1,6):
            cov[i-1][j-1] = pow(.5, (abs(i-j)))
    return cov

```

```

mean = [1,2,3,4,5]
cov = get_cov()

```

```
x = np.random.multivariate_normal(mean, cov, 1000)
```

```
y_epsilon = np.random.randn(1000)
```

```
y = x[:,0] + x[:,1] + np.square((x[:,2] - 2.5)) + y_epsilon
```

```
print(f"First 3 values of X: {x[:3]}")
```

```

First 3 values of X: [[-0.20645625  0.62507688  1.7750502   2.40691521
  5.18134863]
 [ 3.7402364   4.40503372  4.35016723  4.96114578  5.39520319]
 [-0.86926549 -0.28132208  2.10283711  3.41796534  5.90390064]]

```

```
print(f"First 3 values of Y: {y[:3]}")
```

```
First 3 values of Y: [ 0.01941756 12.6972788 -2.12164048]
```

**1.3 1C. [10 Points]** Use the first 400 observations of your data as the training data and the rest as testing data. Predict the Y values using your KNN function with  $k = 5$ . Evaluate the prediction accuracy using mean squared error.

**1.3.1 Answer:**

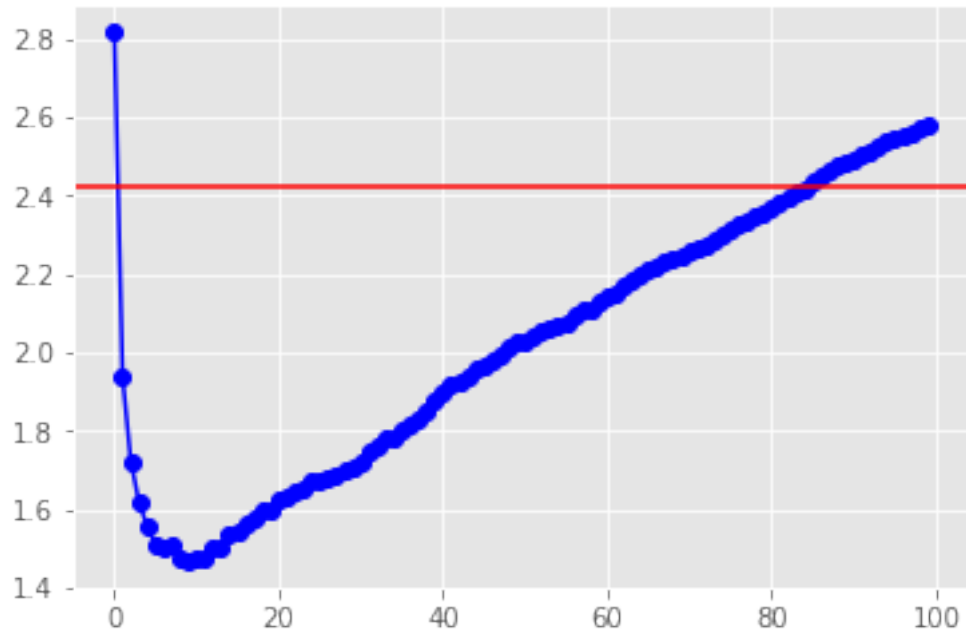
```
def mean_square_error(y_true, y_pred):  
    return np.square(np.subtract(y_true,y_pred)).mean()  
  
X_train, X_test, y_train, y_test = x[:400], x[400:], y[:400], y[400:]  
  
y_pred = myknn(X_test, X_train, y_train, 5)  
  
my_knn_c_mse = mean_square_error(y_test, y_pred)  
  
print(f"MSE: {my_knn_c_mse}")
```

MSE: 1.5558223877014192

**1.4 1D. [15 Points]** Compare the prediction error of a linear model with your KNN model. Consider  $k$  being 1, 2, 3, ..., 9, 10, 15, 20, ..., 95, 100. Demonstrate all results in a single, easily interpretable figure with proper legends.

**1.4.1 Answer:**

```
from sklearn.linear_model import LinearRegression  
import matplotlib.pyplot as plt  
%matplotlib inline  
plt.style.use(['ggplot'])  
  
reg = LinearRegression().fit(X_train, y_train)  
  
lm_y_pred_d = reg.predict(X_test)  
  
lm_mse_d = mean_square_error(y_test, lm_y_pred_d)  
  
k_vals = np.arange(1,101)  
  
knn_msos = np.zeros(100)  
  
for i in range(len(k_vals)):  
    y_pred = myknn(X_test, X_train, y_train, k_vals[i])  
    knn_msos[i] = mean_square_error(y_test, y_pred)  
  
plt.plot(knn_msos, '-o', color='blue')  
plt.axhline(y=lm_mse_d, color='r')  
plt.show()
```



### 1.5 Question 2 [50 Points] Linear Regression through Optimization

1.6 2A. [35 Points] Based on this description, write your own R function `mylm_g(x, y, delta, epsilon, maxitr)` to implement this optimization version of linear regression. The output of this function should be a vector of the estimated beta value.

1.6.1 Answer:

```
def gradient_descent(x, y, b, delta, epsilon, maxitr):
    iterations = 1
    cost_history = []
    m = len(y)

    while True:
        h = x.dot(b)
        loss = h - y

        gradient = x.T.dot(loss) / m

        b_new = b - delta * gradient

        cost = np.sum(abs(b_new - b))
        cost_history.append(cost)

        b = b_new
```

```

        if cost < epsilon:
            break

        if iterations == maxitr:
            break

        iterations +=1

    return b, cost_history

def mylm_g(x, y, delta, epsilon, maxitr):
    beta = np.zeros(x.shape[1])
    newB, _ = gradient_descent(x, y, beta, delta, epsilon, maxitr)
    return newB

```

- 1.7 2B. [15 Points] Test this function on the Boston Housing data from the `mlbench` package. Documentation is provided here if you need a description of the data. We will remove `medv`, `town` and `tract` from the data and use `cmedv` as the outcome. We will use a scaled and centered version of the data for estimation. Please also note that in this case, you do not need the intercept term. And you should compare your result to the `lm()` function on the same data. Experiment on different `maxitr` values to obtain a good solution. However your function should not run more than a few seconds.

#### 1.7.1 Answer:

```

from sklearn.preprocessing import scale

x = pd.read_csv('./boston-data.csv')

x = x.drop([x.columns[0], 'medv', 'town', 'tract'], axis=1)

x = pd.DataFrame(scale(x), index=x.index, columns=x.columns)

y = x.cmedv.to_numpy()

x = x.drop(['cmedv'], axis=1).to_numpy()

# Train params
delta = 0.1
epsilon = 1e-7
max_iterations = 5000

my_beta_2b = mylm_g(x, y, delta, epsilon, max_iterations)

reg = LinearRegression().fit(x, y)

lm_beta_2b = reg.coef_

beta_df = pd.DataFrame({'LinearRegression': lm_beta_2b, 'mylm_g': my_beta_2b})

```

beta\_df

	LinearRegression	mylm_g
0	-0.032316	-0.032317
1	0.030245	0.030245
2	-0.097936	-0.097936
3	0.118273	0.118272
4	0.011390	0.011388
5	0.071312	0.071313
6	-0.199704	-0.199703
7	0.287233	0.287233
8	0.007565	0.007565
9	-0.321039	-0.321039
10	0.290851	0.290846
11	-0.236526	-0.236521
12	-0.206805	-0.206805
13	0.091235	0.091235
14	-0.417973	-0.417973